

# Gale–Shapley (Deferred Acceptance) Algorithm: A Simple Implementation Guide

Francesco Balocco

February 17, 2026

## What problem does Gale–Shapley solve?

You have two equal-sized sets of agents (e.g., **men** and **women**). Each agent has a **ranked list** of the agents on the other side (a strict ordering, from best to worst).

**Goal.** Find a **one-to-one matching** such that there is **no blocking pair**.

**Blocking pair (key concept).** A man  $m$  and a woman  $w$  form a **blocking pair** if:

- $m$  prefers  $w$  to his assigned partner, *and*
- $w$  prefers  $m$  to her assigned partner.

If such a pair exists, the matching is **not stable** because  $m$  and  $w$  would both want to deviate.

## What do you need as input (for coding)?

You need two preference dictionaries:

- `proposer_prefs[p]`: a list of acceptors in descending preference.
- `acceptor_prefs[a]`: a list of proposers in descending preference.

**Example (shape).**

```
proposer_prefs = {
    'W01': ['M04', 'M01', 'M10', ...],
    'W02': ['M10', 'M06', 'M01', ...],
    ...
}
acceptor_prefs = {
    'M01': ['W07', 'W04', 'W01', ...],
    'M02': ['W06', 'W02', 'W08', ...],
    ...
}
```

## How do you get preference lists from data? (the practical step)

In applications (including the speed-dating exercise), you often start from a **score** for each potential pair:

- a predicted probability (e.g.,  $\hat{p} = P(\text{decision} = 1)$ ), or
- any other numeric “utility” score.

To run Gale–Shapley, you only need the **rank order** of those scores:

- For each woman  $w$ : sort all men  $m$  by  $\hat{p}_{w,m}$  from highest to lowest. That is `proposer_prefs[w]` if women propose.
- For each man  $m$ : sort all women  $w$  by  $\hat{p}_{m,w}$  from highest to lowest. That is `acceptor_prefs[m]` in the same run.

**Important detail: ties.** If two scores are exactly equal, break ties deterministically (e.g., by the partner ID). This produces a **strict** ranking, which is what the standard algorithm expects.

## The algorithm (high level intuition)

Gale–Shapley is also called **deferred acceptance**.

- The proposing side proposes to their favorite acceptable option *they have not proposed to yet*.
- The receiving side *tentatively* keeps the best proposal so far (according to their own ranking) and rejects the rest.
- Rejected proposers try again next round with their next best choice.

It stops when no proposer wants (or is able) to make a new proposal.

## The algorithm (exact steps you can implement)

Assume “proposers” propose to “acceptors”.

### Data structures (recommended)

- `free`: queue/list of currently free proposers.
- `next_idx[p]`: which acceptor index proposer  $p$  will propose to next.
- `engaged[a]`: the proposer currently held by acceptor  $a$  (a tentative engagement).
- `acceptor_rank[a][p]`: a fast lookup of how acceptor  $a$  ranks proposer  $p$  (smaller = better).

### Pseudocode

```

build acceptor_rank[a][p] for all acceptors a and proposers p
free = all proposers
next_idx[p] = 0 for all proposers
engaged = {}    # acceptor -> proposer

while free is not empty:
    p = pop one proposer from free
    if next_idx[p] == len(proposer_prefs[p]):
        continue    # p has proposed to everyone

    a = proposer_prefs[p][next_idx[p]]
    next_idx[p] += 1

    if a not in engaged:

```

```

engaged[a] = p
else:
    current = engaged[a]
    if acceptor_rank[a][p] < acceptor_rank[a][current]:
        engaged[a] = p
        push current back into free
    else:
        push p back into free

return match where each proposer p is matched to the acceptor that holds them

```

## How to switch “women propose” vs. “men propose”

Nothing in the code changes except which side you pass in as the proposer dictionaries:

- **Women propose:** proposers = women, acceptors = men.
- **Men propose:** proposers = men, acceptors = women.

**Important output property (one sentence).** The proposing side gets its **best stable outcome** (women-optimal or men-optimal stable matching).

## A tiny example (so you can sanity-check your code)

Women propose. Each woman lists men from best to worst:

```

W1: M1 > M2
W2: M1 > M2
M1: W2 > W1
M2: W1 > W2

```

Round 1: W1→M1, W2→M1. M1 keeps W2 and rejects W1.

Round 2: W1→M2. M2 accepts.

Result: (W2,M1) and (W1,M2). This is stable, and it is best (among stable matchings) for women because women proposed.

## A minimal Python implementation (matches the notebook style)

```

def gale_shapley(proposer_prefs, acceptor_prefs):
    acceptor_rank = {
        a: {p: r for r, p in enumerate(lst)}
        for a, lst in acceptor_prefs.items()
    }
    free = list(proposer_prefs.keys())
    next_idx = {p: 0 for p in proposer_prefs}
    engaged = {} # acceptor -> proposer

    while free:
        p = free.pop(0)

```

```

if next_idx[p] >= len(proposer_prefs[p]):
    continue
a = proposer_prefs[p][next_idx[p]]
next_idx[p] += 1

if a not in engaged:
    engaged[a] = p
    continue

current = engaged[a]
if acceptor_rank[a].get(p, 10**9) < acceptor_rank[a].get(current, 10**9):
    engaged[a] = p
    free.append(current)
else:
    free.append(p)

match = {p: None for p in proposer_prefs}
for a, p in engaged.items():
    match[p] = a
return match

```

## Common implementation pitfalls (and how to avoid them)

- **Slow ranking checks:** do not use `list.index()` inside the main loop for acceptors; pre-compute `acceptor_rank`.
- **Incomplete lists:** the standard proof assumes complete strict rankings. If you drop some partners, decide what “unacceptable” means (usually: they are ranked last or not included at all).
- **Ties:** if two partners have the same score, you must break ties deterministically (e.g., by ID) to get a strict list.
- **Unequal sizes:** the algorithm still runs, but some proposers may end up unmatched (`None`).
- **Interpreting “utility”:** Gale–Shapley uses only *rank order*, not cardinal values. Any monotone transformation of your scores yields the same preferences.

## How to test your code

1. Start with a tiny toy example (2–3 agents per side) where you can compute the outcome by hand.
2. Verify **stability**: check whether any blocking pair exists.
3. Run both regimes (women propose / men propose) and compare outcomes.